Client-Fit & Idea-Tagging Logic

How an idea is scored against a client — today's live engine, and a proposed redesign

Prepared	Scope	Source of truth	Status
18 June 2026	Today's Focus & Solutions Views tagging	mapping.js · scanner.js · data.js	§1–§2 live · §3 proposed

The one-line truth. There is **one** scoring engine — `mapping.js → scoreIdeaForClient(idea, client)`. **Both** Today's Focus and the Views tab score through it. In `scanner.js`, `ideaFit()` simply delegates to `MAPPING.scoreIdeaForClient`, and `clientsForIdea()` ranks the clients it returns. The old standalone scanner formula was deleted when the two engines were unified (commit `007131c`); only a now-unused `CONV_W` export lingers. So — **one engine today, not two.**

All per-axis numbers in the worked examples below were produced by running the actual engine logic (`mapping.js` + the synthetic 24-month sector history generated in `data.js`) against the real seed book for **Aurora**. They are exact, not illustrative.

1 · Current process — Today's Focus mapping Live

Today's Focus surfaces market ideas (earnings prints, structural calls) and, per idea, flags the clients it fits via `MAPPING.flagClients(idea)`. That function scores every client with `scoreIdeaForClient`, keeps those at or above the flag threshold, and returns the top few by fit. The score is a **client-fit** number: how *right* the idea is for *this* book — deliberately separate from the idea's own conviction.

1.1 The four live axes

Each axis returns 0–100 and is combined with **flat (fixed) weights**. After House-view fit was removed from client-fit (moved to the conviction score) in `007131c`, the original five weights were rescaled by `1/0.85` so the surviving four keep their relative balance and sum to **exactly 1.00**. The math uses the exact fractions; the UI shows the rounded values.

Axis (key)	What it measures	Formula (0–100)	Weight (exact)	Shown
Gap fit `gap`	Headroom from the book's current sector allocation up to the strategic peg.	`max(0, (peg − cur) / peg × 100)`	0.20/0.85 = 0.2353	0.24
Affinity fit `holdings`	Recency-weighted (λ=0.94, 24 mo) sector affinity, minus an over-the-peg concentration penalty.	`max(0, ThematicAffinity − Penalty)`	0.25/0.85 = 0.2941	0.29
Mandate & risk `mandate`	Does the idea's vol/beta/structure and goal type suit the mandate?	`0.6·RiskSuitability + 0.4·IntentFit`	0.25/0.85 = 0.2941	0.29
Conc. within sector `concSector`	Herfindahl diversification of the book's in-sector names; *inverted* by default (a concentrated sector wants a new name).	`100 − (1 − HHI)×100`	0.15/0.85 = 0.1765	0.18
Sum of weights			**1.0000**	**1.00**

Axis detail (exact, from `mapping.js`)

- **Gap fit.** `cur` = current % of the book in the idea's sector; `peg` is the strategic comfort limit by mandate class — **growth 25 / income 15 / preservation 10** (one source of truth: `PARAMS.affinity.comfort` via `sectorPeg`, shared with the Affinity penalty). At or over the peg → 0.

- **Affinity fit.** *Thematic affinity* = recency-weighted share (month `t` weight = 0.94^t) of the trailing-24-month months whose allocation sits at or below today's — "has this book been building / sitting at the top of its range". If the book holds the sector but no history is on file → treated as steady → ~100; no exposure → 0. *Penalty* = overshoot above the comfort peg, in pp × 10 (cap 100). Net = max(0, affinity − penalty).

- **Mandate & risk.** *RiskSuitability* grades the idea's vol/beta/structure against the mandate (e.g. income mandate: low-vol → 90, high-beta → 45). *IntentFit* grades the idea's goal type (appreciation / yield / protection) against the mandate's goal. Tradability is **not** on this axis any more.

- **Concentration within sector.** `HHI = Σ(in-sector weighti)²` over weights normalised to sum to 1. Raw diversification = (1−HHI)×100. For the fit blend it is **inverted** (`invertForFit:true`): a concentrated sector scores higher, since a new name diversifies it. No in-sector holding → neutral 50.

1.2 The global tradability gate

Tradability is a **binary {0,1}** multiplier over the *whole* weighted sum — not an axis:

```
Final Client-Fit = Tradability × Σ(axisi · weighti)
```

The model (from `data.js complexityOf`): a **Retail** client cannot trade an **OTC** natural expression (collars, forwards, OTC options). Structured *notes* are packaged securities a Retail client *can* trade; a Professional can trade any expression. The idea's “natural expression” is its first listed structure. If blocked, fit = 0 and the idea is surfaced with the MiFID reason rather than silently dropped.

1.3 The flag threshold

`flagClients` uses `PARAMS.flagMin = 50` and `PARAMS.flagMax = 6`: keep clients with **fit** ≥ **50**, sort by fit descending, take the **top 6**. (Tiers for display: Strong ≥ 66, Good ≥ 48, else Marginal.)

1.4 Worked example — Micron FQ3 × Aurora

Today's Focus headline idea **“Micron FQ3 — the AI-memory print of the quarter”** (`mu-fq3`; sector Technology, bucket Growth, structures `[Phoenix autocall, Zero-cost collar, Buffered note, Direct equity]`) scored for **Aurora** — a **MiFID Retail**, **“Growth, with income needs”** book (parsed to **income mandate** → peg 15%).

Aurora's Technology exposure = Micron 25.8 + NVIDIA 7.2 + SAP 8.4 = **41.4%**. Synthetic 24-mo Tech history runs 6.6–43.5% with today (41.4%) just under last month's 43.5% peak.

Axis	Computation	Score	× Weight	= Contribution
Gap fit	41.4% Tech vs the income peg 15% — far over target, no headroom → `max(0,(15−41.4)/15×100)`.	0	0.2353	0.0
Affinity fit	Thematic affinity **93** (only last month's 43.5% exceeds today's 41.4%) – concentration penalty **100** (26.4pp over the 15% comfort ×10, capped) = max(0, 93–100).	0	0.2941	0.0
Mandate & risk	0.6×RiskSuit **45** (“high-beta — too racy for an income book”) + 0.4×IntentFit **60** (capital-appreciation goal vs an income mandate) = 51.	51	0.2941	15.0
Conc. within sector	3 Tech names (MU/NVDA/SAP), `HHI 0.46` → diversification 54 → inverted **46**.	46	0.1765	8.1
Bracketed weighted sum (round 15.0 + 8.1)		**bracketFit = 23**		
Tradability gate — Retail *can* trade a Phoenix autocall (structured note) → ×1		**Final Fit = 23**		

Why Aurora is *not* flagged (23 < 50) — and why that's correct. Aurora's signature holding *is* Micron (25.8%), yet the idea does not flag her. The engine scores **idea-fit**, not ownership: adding more Technology to a book that is **2.8× over its income-mandate sector peg** wipes out both Gap (0) and Affinity (penalised to 0), and a high-beta growth print only half-suits an income mandate (51). The whole point of Today's Focus is “who should I call about this” — and the honest answer for Aurora is “she already owns too much; this isn't an *add* for her.” The protective angle for her MU concentration is a different conversation (a collar/structure on what she holds), not a fit to buy more.

Contrast — a Today's Focus idea that *does* flag Aurora. The **“Gold — accumulate the pullback”** idea (`gold-pullback`; sector Gold, bucket Protection, natural expression *Physical / ETC*) scores her **82** on the same engine — far above the 50 flag — because she is well *under* peg in Gold with headroom and a clean record there. The arithmetic is identical to the Views worked example in §2.4 (same sector/bucket/structure profile).

2 · Current process — how the Views tab tags ideas to a client Live

The Solutions Views tab uses the **same** `mapping.js` engine. `scanner.js` is a thin facade that preserves the old output keys (`applies / score / reason / gap / secExp / acExp`) while sharing one algorithm with Today's Focus:

```
// scanner.js — the facade
function ideaFit(idea, client) {
return window.MAPPING.scoreIdeaForClient(idea, client); // delegates, no second formula
}

// which clients does a View idea tag?
function clientsForIdea(idea) {
return S().clients
.map(c => ({ client: c, ...ideaFit(idea, c) }))
.filter(x => x.applies) // gate
.sort((a, b) => b.score - a.score); // rank by fit
}
```

The Views gate is `applies`, not the flag. In the engine, `applies = tradable && bracketFit ≥ applyMin`, with **`applyMin = 45`**. So the Views tab tags a client when the idea is tradable for them *and* their pre-gate fit reaches 45 — a slightly looser bar than Today's Focus's **fit ≥ 50, top 6**. Same engine and same per-axis math; only the cut-off and the “top N” cap differ between the two surfaces.

2.1 The honest history

Before unification, `scanner.js` had its **own** standalone formula:

```
// REMOVED in 007131c — the old standalone scanner scorer
score = secExp×1.4 + acExp×0.4 + gap×1.2 + conv
applies = secExp≥6 || acExp≥25 || gap≥6 || (protectionIncome && gap≥4)
```

That math used sector/asset-class exposure, a goal gap, and a conviction weight (`conv` from `CONV_W`), with an `applies` gate that disagreed with the 5-axis engine. Commit `007131c` deleted it and pointed `ideaFit` at `MAPPING.scoreIdeaForClient`. **Removed** the formula · **Dead** `CONV_W = {High:3, "Medium-High":2, Medium:1}` is still exported from `scanner.js` but nothing reads it for scoring any more — it is a vestige, not a live input.

2.2 Worked example — "Gold as a debasement hedge" × Aurora (Views)

The Solutions View idea **"Gold as a debasement hedge"** (`gold-hedge`; sector Gold, bucket Protection, structures `[Physical / ETC, Gold accumulator]` → natural expression *Physical / ETC*, non-complex) scored for **Aurora** (income mandate, Gold peg 15%). Aurora holds one Gold name, **Xetra-Gold ETC at 5.2%**.

Axis	Computation	Score	× Weight	= Contribution
Gap fit	5.2% Gold vs the income peg 15% → `(15–5.2)/15×100` = 65% headroom to the peg.	65	0.2353	15.3
Affinity fit	Thematic affinity **93** (Gold built steadily; today's 5.2% sits near the top of its 24-mo range); 5.2% is *within* the 15% comfort → no penalty.	93	0.2941	27.4

Mandate & risk	0.6×RiskSuit **82** ("moderate-vol yield fits an income mandate") + 0.4×IntentFit **65** (capital-protection goal vs an income mandate) = 75.	75	0.2941	22.1
Conc. within sector	One Gold name → `HHI 1.00` → diversification 0 → inverted **100** (a single concentrated holding — a new gold name diversifies it most).	100	0.1765	17.6
Bracketed weighted sum		**bracketFit = 82**		
Tradability — Retail can hold Physical / ETC (non-complex) → ×1		**Final Fit = 82**		

Does the Views tab tag her? Yes. bracketFit 82 ≥ applyMin 45 and the idea is tradable → `applies = true`, so `clientsForIdea(gold-hedge)` returns Aurora, ranked by her score of 82. (She also clears the stricter Today's Focus flag of 50, which is why the §1 Gold contrast also flags her.)

3 · Proposed redesign Proposed — not built

This section is a design proposal, not the live system. Nothing here is implemented. It reframes the "why does this fit" signal around a **Preference** score (trajectory + house view), demotes the client's **goal** from a weighted axis to a **dial (multiplier)**, and is explicit that **ownership and P&L never add points**.

3.1 The proposed components

Preference (0–100)

```
PREFERENCE = 0.6·Trajectory + 0.4·View
```

Trajectory from a 24-month series `s` (index 0 = most recent):

```
recent = mean(s[0..5]); older = mean(s[18..23]); Δ = recent − older
max(s) = 0 → 0
Δ ≥ +5 → 100
+1 ≤ Δ < +5 → 70 + (Δ−1)/4 × 30
−1 ≤ Δ < +1 → 55
−5 ≤ Δ < −1 → 20 + (Δ+5)/4 × 35
Δ < −5 → max(0, 20 + (Δ+5) × 4)
```

View (the house view on the name/theme):

```
LIKE → 60 + conviction × 4
NEUTRAL → 50
none → 40
AVOID → 10
```

Gap (0–100)

```
GAP = max(0, (peg − current) / peg × 100) peg { growth 25, income 15, preservation 10 }
```

Concentration within sector

```
HHI = Σ(name weight-within-sector)²
raw = (1 − HHI) × 100
contribution = 100 − raw // invert by default
```

Risk suitability

Mandate	high	moderate	low
growth	90	70	45
income	35	75	80
preservation	15	50	90

Goal dial (0.5–1.0) — a multiplier, not an axis

IntentMatch	appreciation	yield	protection
growth	100	45	35
income	45	100	70

preservation	25	70	100

```
dial = 0.5 + 0.5 × (IntentMatch / 100) // ranges 0.5 .. 1.0
```

Tradability gate & composition

```
Tradability = 1 / 0 // Retail + OTC/complex natural expression → 0

Base = 0.40·Preference + 0.20·Gap + 0.15·ConcContribution + 0.25·RiskSuitability
FinalFit = round( Tradability × GoalDial × Base ) flag ≥ 50
```

Goal lives only in the dial. Ownership and P&L never add points — preference is built from the allocation *trajectory* and the *house view*, not from whether the client holds the name or is up/down on it.

3.2 Worked examples (proposed)

Inputs below are illustrative (the redesign is not built, so there is no live data to read); the arithmetic shows the proposed mechanics and lands in the target bands.

Example A — Gold → Amar (growth) → ~43 (drops off via the goal dial)

Amar is a **growth** family office (peg 25) holding Gold at 12% — one Gold name. Gold's goal type is *protection*.

Component	Value	Reasoning
Trajectory	55	Long-held strategic core, roughly steady → Δ in [−1,+1).
View	76	House LIKEs gold, conviction 4 → 60 + 4×4.
Preference	63.4	0.6×55 + 0.4×76 = 33 + 30.4.
Gap	52	(25 − 12)/25 × 100.
Conc. contribution	100	Single gold name → HHI 1 → 100 − 0.
Risk suitability	45	Gold is low-beta; growth×low = 45.
Base	62.0	0.40×63.4 + 0.20×52 + 0.15×100 + 0.25×45 = 25.4 + 10.4 + 15 + 11.25.
Goal dial	0.675	growth×protection IntentMatch 35 → 0.5 + 0.5×0.35.
FinalFit	**≈ 42–43**	**round(1 × 0.675 × 62.0). The goal dial is what knocks a 62-point Base into the low-40s — gold doesn't serve a growth goal.**

Example B — CEG → growth client, built AI-power, down on the year → ~77 **(stays; P&L ignored)**

A **growth** client (peg 25) who has been *building* a power/utilities position (CEG — datacenter power) and is *down on the year*. CEG's goal type is *appreciation*; treat the name as a moderate-beta power play.

Component	Value	Reasoning
Trajectory	100	Allocation built up strongly over 24 mo → Δ ≥ +5. **Note: the trajectory is the *weight* rising, not the price — being down on the year never enters.**
View	76	House LIKEs AI-power, conviction 4 → 60 + 16.
Preference	90.4	0.6×100 + 0.4×76 = 60 + 30.4.
Gap	42	Utilities ~14.5% vs peg 25 → (25–14.5)/25×100.
Conc. contribution	100	Concentrated single power name → new name diversifies.
Risk suitability	70	Moderate-beta power name; growth×moderate = 70.
Base	77.1	0.40×90.4 + 0.20×42 + 0.15×100 + 0.25×70 = 36.16 + 8.4 + 15 + 17.5.
Goal dial	1.00	growth×appreciation IntentMatch 100 → full dial, no penalty.
FinalFit	**77**	**round(1 × 1.00 × 77.1). Stays well above the flag — high preference (trajectory + view) and an on-goal idea; the loss is irrelevant.**

Example C — BABA → holder down hard, no bullish view → ~49/45 **(drops; weak preference)**

A **growth** client holds BABA, *down hard*, with *no bullish house view*. BABA's goal type is *appreciation* (dial 1.0), so the drop is driven purely by a weak **Preference** — not the loss.

Component	Value	Reasoning
Trajectory	37.5	Position faded as it fell → Δ ≈ −3 → 20 + (−3+5)/4×35.
View	40	No bullish house view → “none” = 40.
Preference	38.5	0.6×37.5 + 0.4×40 = 22.5 + 16. **Weak — this is what sinks the score.**
Gap	20	Already heavy in the sector → little headroom.
Conc. contribution	40	Several names in the sector → fairly diversified.
Risk suitability	90	High-beta name; growth×high = 90.
Base	45.7	0.40×38.5 + 0.20×20 + 0.15×40 + 0.25×90 = 15.4 + 4 + 6 + 22.5.
Goal dial	1.00	growth×appreciation = full dial.
FinalFit	**≈ 46**	**round(1 × 45.7). Below the 50 flag → drops off. If the desk actively flags AVOID (View = 10), preference falls further → ~45. Being down hard is *not* a factor.**

The “~49/45” band reflects the two view states: a neutral/absent view lands ~48–49, an active AVOID pulls it to ~45 — either way under the flag, on weak preference rather than P&L.

4 · Current vs proposed — summary

The engines are already unified. Today's Focus and the Views tab have scored through one engine since `007131c`. Unification is *not* a proposed change — it is already shipped. The genuine proposals in §3 are only the two reframings below.

Dimension	Current (live)	Proposed (§3)
Engine	One engine, `scoreIdeaForClient`, shared by Today's Focus & Views Already unified	Same single-engine principle — **no change**
"Why it fits" signal	**Affinity fit** axis: recency-weighted sector affinity – over-peg penalty (rewards having built up / sitting at top of range)	**Preference** = 0.6·Trajectory + 0.4·View — **no ownership reward; P&L removed** (trajectory is the weight's slope; view is the house call)
Client goal	Folded into the **Mandate & risk** axis (IntentFit, 0.4 of a 0.29-weighted axis)	A **dial/multiplier** (0.5–1.0) over the whole Base — an off-goal idea is scaled down, not just lightly penalised
Axes & weights	4 axes: Gap 0.24, Affinity 0.29, Mandate&Risk 0.29, ConcSector 0.18	Base from Preference 0.40, Gap 0.20, ConcContribution 0.15, RiskSuitability 0.25
Tradability	Binary global gate × weighted sum	Binary global gate × dial × Base — **same role**
Flag threshold	Today's Focus fit ≥ 50 (top 6); Views `applies` ≥ 45	FinalFit ≥ 50

In one sentence: the unification work is done; the open proposals are (1) replace the affinity/ownership signal with a **trajectory + house-view Preference** that ignores P&L, and (2) make the client's **goal a multiplier** rather than a weighted term so off-goal ideas are dialled down decisively.

Brokerage Playground — Client-Fit & Idea-Tagging Logic · Prepared 18 June 2026 · Today's Focus data as-of 16 June 2026 · Sources of truth: `mapping.js, scanner.js, data.js` · Engine unification: commit `007131c`. §1–§2 describe the live system; §3 is a proposal, not built.